

Creation of a Redshift Cluster

Screenshots of the configuration of the Redshift cluster that you have created:

Amazon Redshift > Clusters

In my account | From other accounts

► Connect to Redshift clusters

Clusters (1) [Info](#)

[Query data](#) [Actions](#) [Create cluster](#)

<input type="checkbox"/>	Cluster	Cluster namespace	Status	Storage capacity us...	CPU utilization	Snapshots	Notificati...	Tags
<input type="checkbox"/>	dada-etl-redshift-cluster dc2.large 2 nodes 320 GB	c7e15745-40d0-42e8-...	Available	< 1%	35%	2 snapshots		

Amazon Redshift > Clusters > dada-etl-redshift-cluster

dada-etl-redshift-cluster

[Actions](#) [Edit](#) [Add partner integration](#) [Query data](#)

General information

Cluster identifier dada-etl-redshift-cluster	Status Available	Node type dc2.large	Endpoint dada-etl-redshift-cluster.crgbsjlfoqhl.us-east-...
Cluster namespace c7e15745-40d0-42e8-afa1-575d43db7b42	Date created February 06, 2022, 01:02 (UTC+05:30)	Number of nodes 2	JDBC URL jdbc:redshift://dada-etl-redshift-cluster.crgbsjl...
	Storage used 0.02% (0.07 of 320 GB used)	AQUA Not available	ODBC URL Driver={Amazon Redshift (x64)}; Server=dada-...

Amazon Redshift > Configurations > Subnet groups > Subnet group

dada-redshift-etl-cluster

Delete

Modify

Cluster subnet group details

VPC ID	Description	Status
vpc-07254bb1af7ad6d7b	Redshift cluster for ETL	✔ Complete
Attached clusters		
dada-etl-redshift-cluster		

Subnets (1)

Q Search

< 1 >

Availability Zone	Subnet ID	CIDR block
us-east-1e	subnet-0f9ec0435e428840e	10.0.0.0/24

Setting up a database in the Redshift cluster and running queries to create the dimension and fact tables

Queries to create the various dimension and fact tables with appropriate primary and foreign keys:

```
CREATE SCHEMA etl;
```

```
CREATE TABLE etl.loc (
    location_id INTEGER NOT NULL
    ,atm_location VARCHAR(50)
    ,atm_streetname VARCHAR(255)
    ,atm_street_number INTEGER
    ,atm_zipcode INTEGER
    ,atm_lat NUMERIC(10, 3)
    ,atm_lon NUMERIC(10, 3)
    ,PRIMARY KEY (location_id)
);
```

```
CREATE TABLE etl.atm (  
    atm_prim_id INTEGER NOT NULL  
    ,atm_id VARCHAR(20)  
    ,atm_manufacturer VARCHAR(50)  
    ,location_id INTEGER  
    ,PRIMARY KEY (atm_prim_id)  
    ,FOREIGN KEY (location_id) REFERENCES etl.loc(location_id)  
    );
```

```
CREATE TABLE etl.DATE (  
    year INTEGER  
    ,month VARCHAR(20)  
    ,day INTEGER  
    ,hour INTEGER  
    ,weekday VARCHAR(20)  
    ,full_date_time TIMESTAMP  
    ,date_id INTEGER  
    ,PRIMARY KEY (date_id)  
    );
```

```
CREATE TABLE etl.card (  
    card_type VARCHAR(23)  
    ,card_type_id INTEGER  
    ,PRIMARY KEY (card_type_id)  
    );
```

```
CREATE TABLE etl.FACT_ATM_TRANS (
    trans_id BIGINT NOT NULL
    ,atm_prim_id INTEGER
    ,location_id INTEGER
    ,date_id INTEGER
    ,card_type_id INTEGER
    ,atm_status VARCHAR(20)
    ,currency VARCHAR(10)
    ,service VARCHAR(20)
    ,transaction_amount INTEGER
    ,message_code VARCHAR(255)
    ,message_text VARCHAR(255)
    ,rain_3h NUMERIC(10, 3)
    ,clouds_all INTEGER
    ,weather_id INTEGER
    ,weather_main VARCHAR(50)
    ,weather_description VARCHAR(255)
    ,PRIMARY KEY (trans_id)
    ,FOREIGN KEY (atm_prim_id) REFERENCES etl.atm(atm_prim_id)
    ,FOREIGN KEY (location_id) REFERENCES etl.loc(location_id)
    ,FOREIGN KEY (date_id) REFERENCES etl.DATE (date_id)
    ,FOREIGN KEY (card_type_id) REFERENCES etl.card(card_type_id)
);
```

Select database [Info](#)
To view schemas, select a database.

dev ▼

Select schema [Info](#)
To view tables, select a schema.

etl ▼

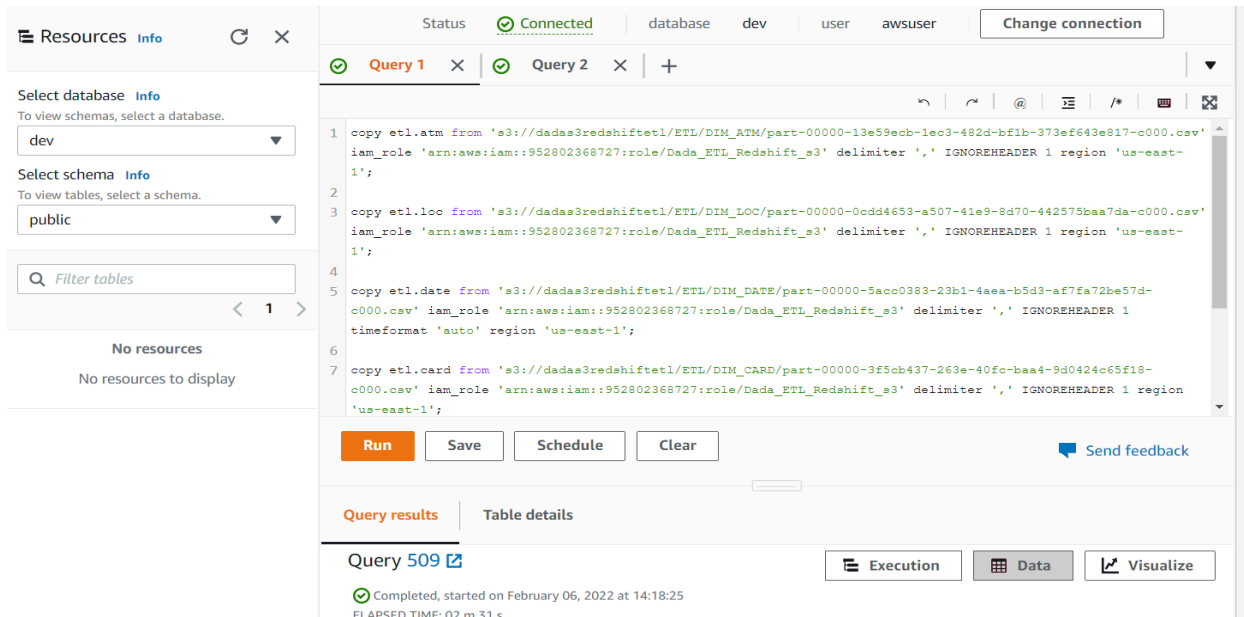
🔍 Filter tables

< 1 >

▶ atm_pkey	...
▶ card_pkey	...
▶ date_pkey	...
▶ fact_atm_trans_pkey	...
▶ loc_pkey	...
▶ atm	...
▶ card	...
▶ date	...
▶ fact_atm_trans	...
▶ loc	...

Loading data into a Redshift cluster from Amazon S3 bucket

Queries to copy the data from S3 buckets to the Redshift cluster in the appropriate tables



The screenshot shows the AWS Redshift console interface. On the left, there's a sidebar with 'Resources' and 'Info' tabs. Below 'Resources', it says 'No resources' and 'No resources to display'. The main area shows a list of queries. The first query is selected, and its SQL code is displayed in the editor. The code is as follows:

```
1 copy etl.atm from 's3://dadas3redshiftetl/ETL/DIM_ATM/part-00000-13e59ecb-1ec3-482d-bf1b-373ef643e817-c000.csv'
iam_role 'arn:aws:iam::952802368727:role/Dada_ETL_Redshift_s3' delimiter ',' IGNOREHEADER 1 region 'us-east-1';
2
3 copy etl.loc from 's3://dadas3redshiftetl/ETL/DIM_LOC/part-00000-0cdd4653-a507-41e9-8d70-442575baa7da-c000.csv'
iam_role 'arn:aws:iam::952802368727:role/Dada_ETL_Redshift_s3' delimiter ',' IGNOREHEADER 1 region 'us-east-1';
4
5 copy etl.date from 's3://dadas3redshiftetl/ETL/DIM_DATE/part-00000-5acc0383-23b1-4aea-b5d3-af7fa72be57d-c000.csv' iam_role 'arn:aws:iam::952802368727:role/Dada_ETL_Redshift_s3' delimiter ',' IGNOREHEADER 1
timeformat 'auto' region 'us-east-1';
6
7 copy etl.card from 's3://dadas3redshiftetl/ETL/DIM_CARD/part-00000-3f5cb437-263e-40fc-baa4-9d0424c65f18-c000.csv' iam_role 'arn:aws:iam::952802368727:role/Dada_ETL_Redshift_s3' delimiter ',' IGNOREHEADER 1 region
'us-east-1';
```

Below the queries, there are buttons for 'Run', 'Save', 'Schedule', and 'Clear'. At the bottom, there's a section for 'Query results' and 'Table details'. The 'Query results' section shows 'Query 509' and its execution status: 'Completed, started on February 06, 2022 at 14:18:25' with an 'ELAPSED TIME: 02 m 31 s'.

copy etl.atm from 's3://dadas3redshiftetl/ETL/DIM_ATM/part-00000-13e59ecb-1ec3-482d-bf1b-373ef643e817-c000.csv' iam_role 'arn:aws:iam::952802368727:role/Dada_ETL_Redshift_s3' delimiter ',' IGNOREHEADER 1 region 'us-east-1';

copy etl.loc from 's3://dadas3redshiftetl/ETL/DIM_LOC/part-00000-0cdd4653-a507-41e9-8d70-442575baa7da-c000.csv' iam_role 'arn:aws:iam::952802368727:role/Dada_ETL_Redshift_s3' delimiter ',' IGNOREHEADER 1 region 'us-east-1';

copy etl.date from 's3://dadas3redshiftetl/ETL/DIM_DATE/part-00000-5acc0383-23b1-4aea-b5d3-af7fa72be57d-c000.csv' iam_role 'arn:aws:iam::952802368727:role/Dada_ETL_Redshift_s3' delimiter ',' IGNOREHEADER 1 timeformat 'auto' region 'us-east-1';

copy etl.card from 's3://dadas3redshiftetl/ETL/DIM_CARD/part-00000-3f5cb437-263e-40fc-baa4-9d0424c65f18-c000.csv' iam_role 'arn:aws:iam::952802368727:role/Dada_ETL_Redshift_s3' delimiter ',' IGNOREHEADER 1 region 'us-east-1';

copy etl.FACT_ATM_TRANS from 's3://dadas3redshiftetl/ETL/FACT_ATM_TRANS/part-00000-b63eac25-f44b-4363-a568-5bca6f00f734-c000.csv' iam_role 'arn:aws:iam::952802368727:role/Dada_ETL_Redshift_s3' delimiter ',' IGNOREHEADER 1 region 'us-east-1' TRUNCATECOLUMNS CSV;